

# Comparison of Graph Processing Systems

Tuesday, 20<sup>th</sup> October 2020

# Motivation and Goal

- graph processing becomes increasingly important in academic and industrial environments
  - many problems modeled with graphs, e. g., machine learning and data mining
  - many business models are based on graphs, e. g., viral marketing or Google's search engine
  - graph sizes increase to several billion edges
- performance, parallelism and distribution of graph algorithms becomes more important

*Main Goal: Comparison of five graph processing systems  
in their performance on different graphs and algorithms.*

# Overview

## 1. Preliminaries

- Basics
- Computation Styles
- Hugepages

## 2. Frameworks

## 3. Evaluation

- Research vs. Production Case
- Results

## 4. Conclusion and Outlook


# Preliminaries

## Graphs

A *weighted, directed graph* is the tuple  $G = (V, E, w)$  where the *vertex set* is  $V \subseteq \mathbb{N}$  and the  $E$  is the *edge set* with

$$E \subseteq \{(x, y) \mid x, y \in V, x \neq y\}$$

and  $w : E \rightarrow \mathbb{R}$  is a mapping of edge to a weight.




## Algorithms

**Single-Source Shortest-Paths (SSSP):** find the shortest path from a starting vertex to every other vertex

**Breadth-first search (BFS):** find a node outgoing from a starting vertex, by increasing maximum hop count step-wise


**PageRank (PR):** link analysis algorithm; weighs vertices, measuring their relative importance

# Push Style



- reads active vertex, writes neighborhood
- more efficient, if only few active vertices at the same time
- more efficient, if neighborhoods of active vertices do not overlap

# Pull Style



- reads neighborhood, writes active vertex
- only one write and many read operations
- less synchronization in parallel implementations needed
  - more efficient, if many vertices active at the same time

# Hugepages

- most systems use virtual memory management
  - represents an abstraction to hardware memory
  - virtual memory is organized in pages
  - translations of virtual memory to physical memory are cached, because every translation takes time
- typically, memory pages are 4 KiB in size
- **hugepages** can be several MiB in size → reduce number of cache misses
- especially noticeable in very memory intensive applications

# Frameworks

Framework	Version	NUMA	Dist.	Features	Notes
■ <b>Galois</b>	29.06.2020	✓	(✓)	general purpose library designed for parallel programming, Hugepage support	distributed using Gluon
■ <b>Gemini</b>	02.11.2016	✓	✓	distributed message-based approach from scratch	version contains bugs that had to be fixed
■ <b>Giraph</b>	08.05.2020	X	✓	built on Apache Hadoop	BFS is not natively supported
■ <b>Ligra</b>	14.08.2019	✓	X	dynamically switches between push and pull style	
■ <b>Polymer</b>	28.08.2018	✓	X	optimizes data layout and memory access strategies	

# Evaluation

## Machines

vsflash1-5,

- 96 cores, of which 48 virtual
- 256 GB of RAM each<sup>1</sup>
- Ubuntu 18.04.2 LTS

## Measurements


- **execution time:** time from start to finish of the console command
- **calculation time:** time the framework actually executed the algorithm
- executed each test case 10 times

---

<sup>1</sup>one machine only 128 GB

## Graphs

Both rMat graphs are synthetic, others are real-world data sets; Flickr: 24MB, rMat28: 76GB





# Production Case

- running system: multiple calculations on a single graph
  - graph data stays loaded between calculations
- short calculation times should be preferred
- Not main focus of this presentation!<sup>2</sup>


# Research Case

- individual calculation cases: possibly new graph for each calculation
  - frequently changing algorithm
- framework should be relatively fast on different algorithms
- overall small execution times should be preferred


---

<sup>2</sup>see paper for details

# Production Case Distributed



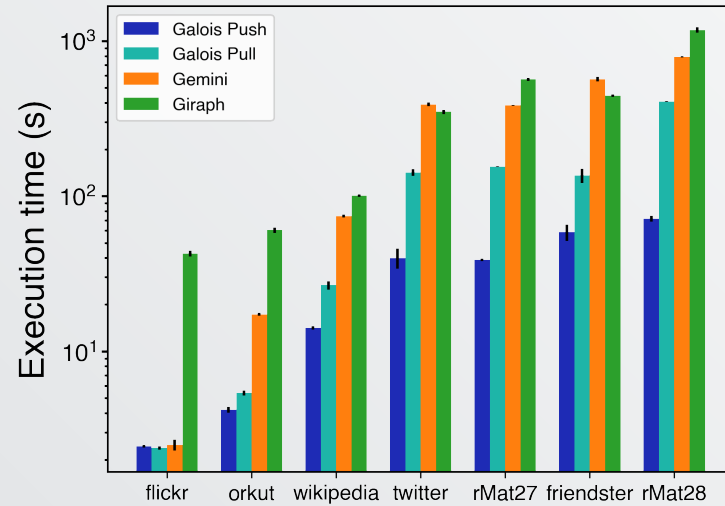
(a) SSSP



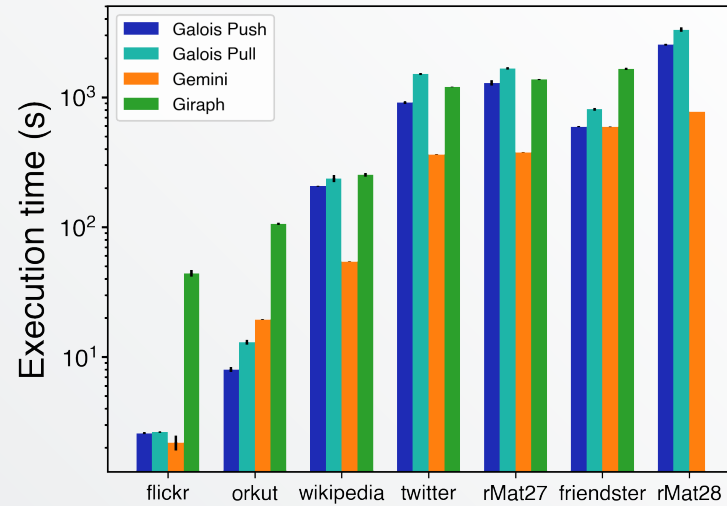
(b) PR

- Giraph is fastest on SSSP and BFS on the real world graphs
- Giraph has problems with synthetic graphs
- Gemini is fastest on PR, with Giraph on second place

# Research Case Distributed




(a) SSSP



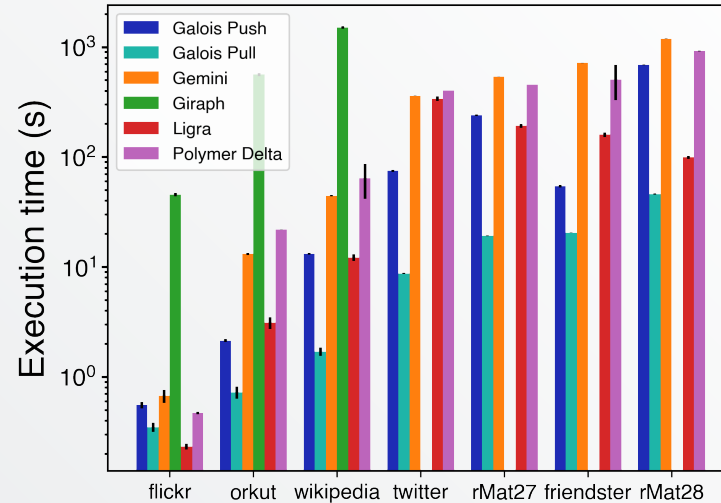
(b) PR

- Galois Push is faster than Pull in all cases
- Both Galois implementations fastest on SSSP or BFS
- Gemini is fastest on PR in almost all cases

# Research Case Single Node



(a) SSSP



(b) PR

- Giraph is either slowest or requires too much RAM (>256 GB)
- Galois is fastest in almost all cases, second fastest is Ligra
- Gemini and Polymer are comparably slow

# Galois With Hugepages

Graph	Calc Time (s)		Exec Time (s)	
	w/o	w/	w/o	w/
flickr	0.01	<b>0.01</b>	0.3	<b>0.2</b>
orkut	0.10	<b>0.02</b>	0.8	<b>0.5</b>
wikipedia	0.38	<b>0.11</b>	1.8	<b>1.1</b>
twitter	2.47	<b>0.94</b>	10.8	<b>5.1</b>
rMat27	4.50	<b>1.39</b>	16.0	<b>6.4</b>
friendster	4.70	<b>1.78</b>	14.4	<b>7.5</b>
rMat28	9.77	<b>3.34</b>	27.8	<b>13.1</b>


(a) SSSP

Graph	Calc Time (s)		Exec Time (s)	
	w/o	w/	w/o	w/
flickr	0.01	<b>0.01</b>	0.3	<b>0.2</b>
orkut	0.06	<b>0.02</b>	0.7	<b>0.6</b>
wikipedia	0.17	<b>0.03</b>	1.7	<b>1.4</b>
twitter	0.77	<b>0.11</b>	<b>8.7</b>	9.3
rMat27	0.65	<b>0.13</b>	19.2	<b>8.1</b>
friendster	1.01	<b>0.14</b>	20.4	<b>13.1</b>
rMat28	1.15	<b>0.24</b>	46.0	<b>16.4</b>


(b) PR Pull

- Hugepages reduce both calculation and execution time on all algorithms
- Execution times can be up to 3× shorter

# Multithreaded Speedup of Galois



(a) SSSP with Hugepages



(b) PR Pull with Hugepages

- Speedups can be significant, with and without hugepages
- Speedup of PR not to the same degree as on SSSP (2.5× vs. 15×)

# Conclusion and Outlook

Generally: 1) performance highly dependent on the framework, algorithm and data set  
2) single node almost always preferable, as long as RAM is sufficient

## Production Case

- Giraph is very fast on distributed systems (especially SSSP and BFS)
- Gemini is fast for distributed PR
- Gemini and Ligra are good options for single node

## Research Case

- Galois is fastest in almost all cases; further improvements with hugepages possible

## Outlook

- incorporate new frameworks and new algorithms
- explore range of settings and other implementations
- repeat similar tests in the future: frameworks are updated and new ones are introduced

# Additional Data






Figure 6: Average times for SSSP on a single computation node



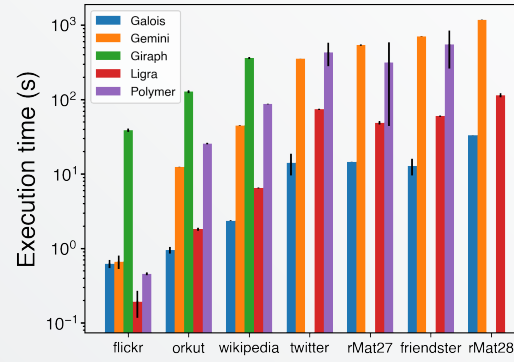



Figure 7: Average times for SSSP on the distributed cluster



(a) Calculation time

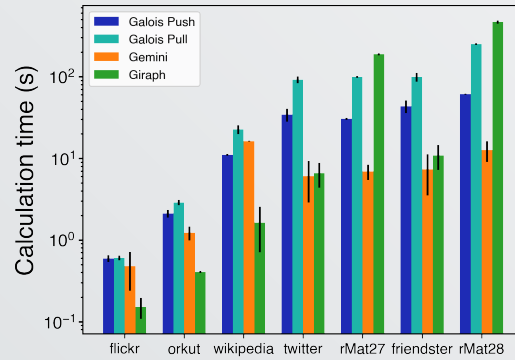


(b) Execution time

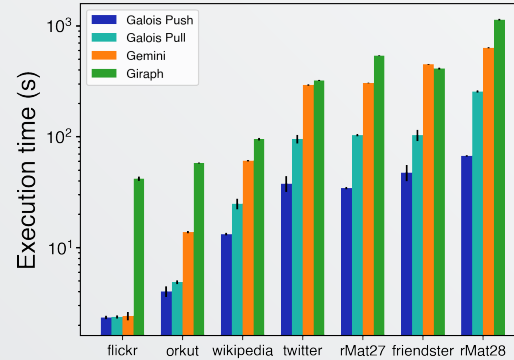


(c) Overhead time

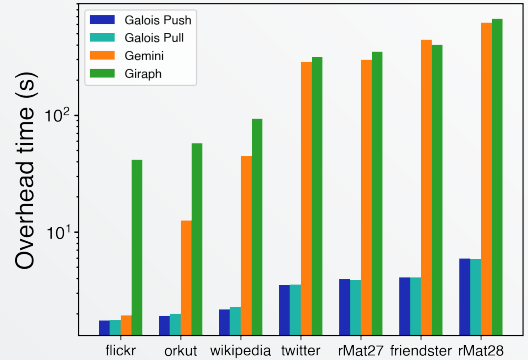
Figure 8: Average times for BFS on a single computation node



(a) Calculation time



(b) Execution time



(c) Overhead

Figure 9: Average times for BFS on the distributed cluster

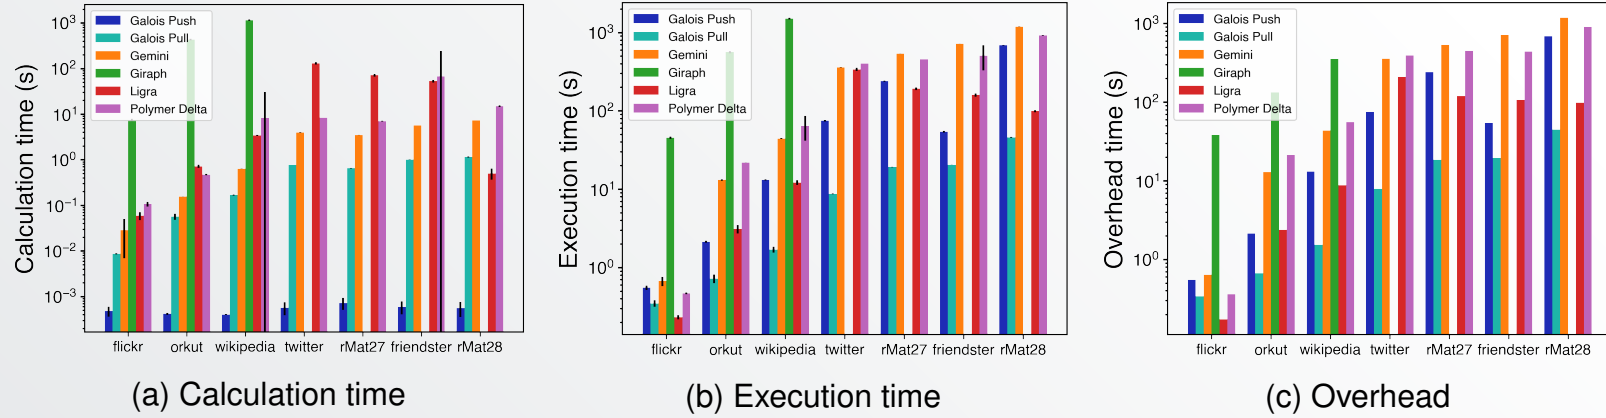


Figure 10: Average times for PR on a single computation node



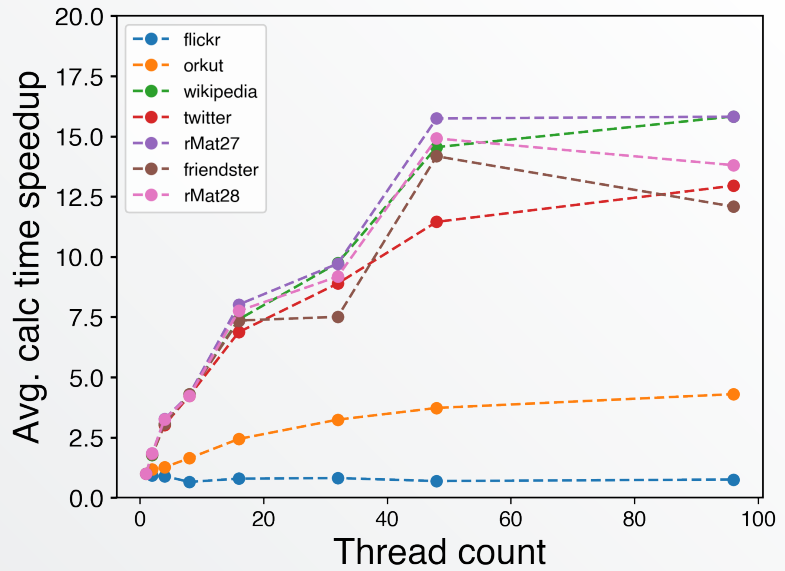


Figure 11: Average times for PR on the distributed cluster




(a) without hugepages




(b) with hugepages

Figure 12: Calculation time speedups on SSSP



(a) without hugepages



(b) with hugepages

Figure 13: Calculation time speedups on PR Pull